

# Short-Term Forecasting of Household Appliance Electricity Demand

---

## 1. Introduction

The push toward heat pumps, solar power, and EVs has made single-household electricity demand much more volatile. Forecasting at the individual home level is considerably harder unlike grid substations that average out noise, household data reflects every random decision a resident makes.

Using the 137-day Belgian low-energy house dataset from Candanedo et al. (2017), this work evaluates 24-hour-ahead appliance demand across four model types: standard benchmarks, SARIMAX, direct gradient boosting, and a zero-shot foundation model.

### Main Takeaways:

**Minimal Real-World Advantage:** Gradient boosting lowered MASE from 0.813 (weekly persistence) to 0.708 a 12.9% improvement. Yet, a Diebold–Mariano test ( $p = 0.34$ ) proved the difference was statistically meaningless.

**Surface-Level Learning:** Over 90% of feature importance was driven purely by calendar variables. The model did not learn real usage habits; it just tracked time-of-day averages.

**Methodological Rigor:** Many published papers rely on contemporaneous sensor data that is not accessible at prediction time. This study avoids that trap by adhering strictly to realistic deployment boundaries.

## 2. Data and preprocessing

The dataset comprises 19,735 complete ten-minute records from 11 January to 27 May 2016, across 28 columns (Candanedo, Feldheim and Deramaix, 2017).

To keep the seasonal SARIMAX estimate practical, we cut the daily period from 144 steps down to 24 by resampling the sequence to hourly means, resulting in 3,290 gap-free observations. While this aggregation irons out brief appliance switching spikes, it converts the target metric to the average 10-minute consumption rate per hour rather than total energy. Environmental sensor feeds (indoor/outdoor temperatures and humidity) were averaged hourly as well.

## 3. Exploratory analysis and stationarity

### 3.1 Distributional characteristics

| Statistic                | Value  |
|--------------------------|--------|
| Mean                     | 97.78  |
| Median                   | 63.33  |
| Standard deviation       | 81.21  |
| Minimum                  | 28.33  |
| Maximum                  | 608.33 |
| Skewness                 | 2.39   |
| Excess kurtosis          | 6.33   |
| 95th percentile / median | 4.37   |

### 3.2 Seasonal structure

Daily electricity consumption ranges from a low of 48.2 Wh at 03:00 to a peak hours of 191.8 Wh at 18:00 a peak-to-trough ratio of 3.98. This sharp contrast aligns with a typical occupied home, where cooking and general appliance use spike right after the working day.

Seasonal strength was quantified using the STL-based measure of Wang, Smith and Hyndman (2006): daily strength (period 24) is 0.318 and weekly strength (period 168) is 0.393.




Figure 1: hourly series, mean profile by hour of day, and mean by day of week.

Both values are moderate at best.

### 3.3 Stationarity

ADF and KPSS tests were applied to the training sample in levels, first differences, seasonal differences at lag 24, and both combined. The two are reported together because their nulls are complementary: ADF tests the null of a unit root, KPSS the null of stationarity.

| Series                   | ADF stat | ADF $p$ | KPSS stat | KPSS $p$ | Reading    |
|--------------------------|----------|---------|-----------|----------|------------|
| Level                    | -8.761   | <0.001  | 0.061     | >0.10    | Stationary |
| First difference         | -15.948  | <0.001  | 0.042     | >0.10    | Stationary |
| Seasonal difference (24) | -12.672  | <0.001  | 0.015     | >0.10    | Stationary |
| Both differences         | -17.445  | <0.001  | 0.061     | >0.10    | Stationary |

## 4. Forecasting design

### 4.1 Rolling-origin protocol

Real-world forecasting relies on daily updates as new data arrives, making a single, static 336-step evaluation unrealistic for a 24-hour-ahead task. Rolling-origin evaluation is the standard approach for assessing out-of-sample forecast accuracy (Tashman, 2000). To match actual operations, the primary experiment sets 14 origins spaced 24 hours apart. At each origin, models process all available history, including test points released from previous windows, to issue a 24-step forecast. Combining these 14 blocks creates a full 336-step prediction series on the test index, giving every model the exact same information set at each step.

### 4.2 Metrics

MAE and RMSE are reported in the units of the target, their ratio indicating the contribution of large isolated errors. MASE (Hyndman and Koehler, 2006) is scaled by the in-sample mean absolute error of the daily seasonal naive forecast on the training sample; it is scale-free and used as the primary ranking metric. Bias is the mean of prediction minus actual, so positive values indicate over-forecasting. Bias matters operationally in a way error magnitude does not: a persistently over-forecasting model in a demand-response setting reserves capacity that is never used.

### 4.3 Forecast origins and covariate availability

Covariates were partitioned based on real-time operational availability. Calendar features (hour, day of week, cyclical encodings) are known indefinitely and used unshifted. Indoor sensor data (temperature and humidity across nine locations) was lagged by the forecast horizon both to prevent data leakage and because indoor climate responds dynamically to occupancy and appliance use. Outdoor weather was evaluated under two setups: *conditional* (using realized test-set values) and *operational* (persisting the last known observation). Finally, lights, rv1, and rv2 were completely excluded.

### 4.4 Leakage controls

Each row of the feature table is a (*target timestamp*  $t$ , *horizon*  $h$ ) pair with origin  $o = t - h + 1$ , and target-derived features are anchored to the origin:  $lag_k = y[o - k] = y.shift(h - 1 + k)$ , with rolling windows spanning

$[o - w, o - l]$ . Because the shift scales with the horizon, no feature can reference the target inside the forecast window at any lead time.

## 5. Benchmark models

Comparing complex models against simple benchmarks is central to the forecasting literature (Makridakis, Spiliotis and Assimakopoulos, 2018). Five benchmarks were evaluated under the same protocol as every other model: the training mean; the naive forecast; daily and weekly seasonal naive rules at lags 24 and 168; and drift (Hyndman and Athanasopoulos, 2021).

| Benchmark             | MAE   | RMSE   | MASE  | Bias   |
|-----------------------|-------|--------|-------|--------|
| Weekly seasonal naive | 43.46 | 81.41  | 0.813 | −13.16 |
| Daily seasonal naive  | 48.31 | 85.57  | 0.904 | 1.75   |
| Mean                  | 50.26 | 74.94  | 0.941 | −3.29  |
| Naive                 | 85.55 | 110.39 | 1.601 | 50.98  |
| Drift                 | 85.80 | 110.68 | 1.606 | 51.37  |

**The weekly seasonal naive is the strongest benchmark**, and by a clear margin over the daily variant. This follows directly from the seasonal strength results in Section 3.2, where weekly strength (0.393) exceeded daily (0.318), and it reveals something specific about the load. The household’s routine is day-specific: Monday and Friday differ systematically from Tuesday, and a rule that looks back one week preserves that distinction while a rule that looks back one day destroys it.

## 6. SARIMAX

### 6.1 Specification

A SARIMAX(1,0,1)(1,0,1)24 model with a constant was estimated by maximum likelihood on the training sample, with the five weather channels as exogenous regressors. Differencing orders follow from Section 3.3. Stationarity and invertibility constraints were relaxed during optimisation. Estimation took 223 seconds and produced AIC 32,506.5 and BIC 32,572.3. SARIMAX was selected because the stationarity tests in Section 3.3 identified a series that is stationary in levels with pronounced daily seasonality, which a seasonal ARMA specification models directly, while the exogenous term admits the weather channels.

### 6.2 Rolling-origin implementation

Parameters were estimated once on the training sample. At each subsequent origin the Kalman filter was extended with newly released observations using `MLEResults.append (refit=False)`, so the state conditions on all data up to the origin while coefficients remain at training values.

### 6.3 Conditional and operational variants

| Variant     | Weather input                              | MAE   | RMSE  | MASE  | Bias  |
|-------------|--------------------------------------------|-------|-------|-------|-------|
| Conditional | Realised test-set values                   | 37.81 | 66.42 | 0.708 | −5.25 |
| Operational | Persisted from last pre-origin observation | 38.51 | 66.38 | 0.721 | −3.83 |

The gap is 0.013 MASE, a 1.9 per cent relative degradation, and a Diebold–Mariano test returns  $DM = -1.49$ ,  $p = 0.136$  not significant at the 5 per cent level.

Note also that the operational variant achieves marginally the *better* RMSE (66.38 against 66.42) despite the worse MAE.

### 6.4 Comparison with the strongest benchmark

Conditional SARIMAX improves on the weekly seasonal naive by 0.105 MASE, from 0.813 to 0.708, a relative improvement of 12.9 per cent. A Diebold–Mariano test (Diebold and Mariano, 1995) on the absolute-error differential returns  $DM = -0.96$ ,  $p = 0.34.7$ . Feature-based model

### 7.1 Formulation

The design matrix contained 74,652 rows and 47 features, of which 66,588 fell in the training period. Gradient boosting was chosen over a linear model because the relationship between hour of day and consumption is

strongly non-linear, and over a recursive formulation because supplying the horizon as a feature avoids the error accumulation and train/apply distribution shift that recursive prediction introduces across 24 steps.

## 7.2 Hyperparameters

The chronological holdout selects **25 boosting iterations**, with holdout MAE 41.21.

| Iterations | Holdout MAE |
|------------|-------------|
| 1          | 47.26       |
| 25         | 41.21       |
| 50         | 43.23       |
| 100        | 44.18       |
| 200        | 46.68       |
| 400        | 48.95       |
| 1500       | 50.68       |

Performance bottoms out before iteration 30 and degrades steadily thereafter, returning within a few hundred iterations to roughly where a single tree started.

## 7.3 Feature importance

Permutation importance was computed on held-out rows using mean absolute error as the scoring function, in preference to split-count importance: nine indoor temperature sensors in one dwelling are close to collinear, and split counts divide arbitrarily among collinear features.

| Feature      | Importance | SD    |
|--------------|------------|-------|
| hour         | 11.252     | 0.545 |
| hour_sin     | 2.304      | 0.542 |
| dow_sin      | 2.058      | 0.480 |
| RH_5_origin  | 0.631      | 0.220 |
| roll_std_168 | 0.598      | 0.149 |
| RH_3_origin  | 0.218      | 0.183 |
| dayofweek    | 0.208      | 0.460 |
| hour_cos     | 0.168      | 0.378 |
| RH_4_origin  | 0.089      | 0.034 |
| roll_std_24  | 0.083      | 0.052 |

## 7.4 Performance

The feature model achieves MAE 37.97, RMSE 65.17, MASE 0.711, Bias  $-5.23$ . It ranks second overall by MASE, a hair behind conditional SARIMAX (0.708), and **first on RMSE** 65.17 against 66.42. A Diebold–Mariano test against SARIMAX returns  $DM = -0.07$ ,  $p = 0.94$ ; the two are as close to indistinguishable as the test can express. Against the weekly seasonal naive,  $p = 0.26$ .

## 8. Foundation model

### 8.1 Configuration

Chronos-T5 (small variant, approximately 46 million parameters) was applied zero-shot. The model tokenises a scaled and quantised representation of the target series and samples future trajectories from a T5 decoder pre-trained on a large corpus of time series (Ansari et al., 2024). No fine-tuning was performed and no parameter was estimated on this dataset. Chronos was included to test whether pre-training on unrelated time series transfers to a dwelling the model has never observed, a question none of the other three model classes addresses.

**The model is univariate.** It receives only the target history no calendar features, no weather, no indoor sensors.

## 8.2 Results

| Metric | Chronos      | SARIMAX (cond.) | Feature model | Weekly seasonal naive |
|--------|--------------|-----------------|---------------|-----------------------|
| MAE    | <b>36.23</b> | 37.85           | 37.97         | 43.46                 |
| RMSE   | 74.40        | 66.41           | <b>65.17</b>  | 81.41                 |
| MASE   | <b>0.678</b> | 0.708           | 0.709         | 0.813                 |
| Bias   | −29.87       | −5.14           | −4.00         | −13.16                |

Read on MASE alone, the foundation model wins: 0.678 against 0.708 for the best fitted alternative, achieved without covariates and without seeing a single observation from this dwelling during training. That would be a striking result.

## 8.3 The point summary drives the ranking

The bias of −29.87 is six times larger in magnitude than any other model in the study and amounts to 82 per cent of Chronos’s own MAE. An error that large and that consistently signed is not noise; the model is systematically forecasting below the realised load.

## 8.4 Sampling variance and interval calibration

Because Chronos relies on stochastic sampling, its final predictions naturally vary from run to run. Evaluating the model across three random seeds produced MASE scores of 0.6704, 0.6971, and 0.6869 showing a spread of 0.0267 due purely to random variation.

A repeat test confirmed that this ~0.027 variation happens consistently. Crucially, that noise level is practically the same size as the performance gap between Chronos and conditional SARIMAX (0.678 vs. 0.708). Because the gap is so close to the random noise margin, the evidence does not support that Chronos is actually the better model.

## 8.5 Assessment

The honest summary is that the foundation model performs **comparably** to the fitted models rather than better than them. Its nominal advantage on MASE dissolves under the point.

## 9. Results and error analysis

### 9.1 Aggregate performance

| Model                 | MAE          | RMSE         | MASE         | Bias   |
|-----------------------|--------------|--------------|--------------|--------|
| Chronos (zero-shot)   | <b>36.23</b> | 74.40        | <b>0.678</b> | −29.87 |
| SARIMAX (conditional) | 37.85        | 66.41        | 0.708        | −5.14  |
| Feature model         | 37.97        | <b>65.17</b> | 0.711        | −4.00  |
| SARIMAX (operational) | 38.43        | 66.35        | 0.719        | −3.93  |
| Weekly seasonal naive | 43.46        | 81.41        | 0.813        | −13.16 |
| Daily seasonal naive  | 48.31        | 85.57        | 0.904        | 1.75   |
| Mean                  | 50.26        | 74.94        | 0.941        | −3.29  |
| Naive                 | 85.55        | 110.39       | 1.601        | 50.98  |
| Drift                 | 85.80        | 110.68       | 1.606        | 51.37  |

Diebold–Mariano tests (Diebold and Mariano, 1995) against the leading model, using a Newey–West variance estimator (Newey and West, 1987) at 23 lags:

Three observations follow.

**The metric chosen determines the winner.** Chronos leads on MAE and MASE and comes seventh of nine on RMSE; the feature model leads on RMSE and comes third on MASE; the flat mean forecast is sixth on MAE but fourth on RMSE.




Figure 2: forecasts against actuals, first 72 test hours; dotted lines mark origins.

## 9.2 Error growth with lead time




Figure 3: mean absolute error against lead time.

**These curves must be read with the confound identified in Section 4.1.** Because origins fall at 19:00 and are spaced at exactly 24 hours, lead time  $h$  corresponds to clock hour  $(19 + h) \bmod 24$ . Lead 6 is 01:00, lead 12 is 07:00, lead 18 is 13:00, lead 24 is 19:00.

**The advantage erodes fastest for the model that has it.** Between lead 1 and lead 24 SARIMAX degrades by a factor of 3.07, the feature model by 1.87, and the seasonal naive rules by 1.71 to 2.07.

## 9.3 Residual diagnostics




Figure 4: residual series, distribution and autocorrelation, conditional SARIMAX.

Residuals of the leading model (conditional SARIMAX, prediction minus actual) fail every diagnostic.

| Diagnostic               | Value                                | Reading                          |
|--------------------------|--------------------------------------|----------------------------------|
| Ljung–Box, 24 lags       | $Q = 204.8, p = 1.3 \times 10^{-30}$ | Strong residual autocorrelation  |
| Ljung–Box, 48 lags       | $Q = 261.7, p = 3.4 \times 10^{-31}$ | Persists                         |
| Residual ACF, lag 1      | 0.564                                | Far outside the $\pm 0.107$ band |
| Residual ACF, lag 24     | 0.156                                | Outside the band                 |
| Skewness                 | -2.68                                | Strong left skew                 |
| Excess kurtosis          | 8.68                                 | Heavy tails                      |
| Corr( residual , actual) | 0.841                                | Severe heteroskedasticity        |

**Residual autocorrelation at lag 1 of 0.564 is very large.** Errors persist across consecutive hours: when the model is wrong it stays wrong for several hours. This is not a specification failure that respecifying the ARMA

order would fix, because the residuals are computed from 24-step-ahead forecasts, not one-step-ahead in-sample residuals.

**The left skew of  $-2.68$  is diagnostic of the central difficulty.** Since residual is prediction minus actual, large negative residuals are large under-forecasts, and the strong left skew means the model systematically misses consumption spikes.

## 9.4 Long-horizon behaviour

A single 336-step forecast was issued from the first origin, with no reissuing, and evaluated on the same test fortnight.

| Model                 | MAE    | RMSE   | MASE  | Rolling-origin MASE |
|-----------------------|--------|--------|-------|---------------------|
| SARIMAX (conditional) | 37.91  | 66.34  | 0.710 | 0.708               |
| Weekly seasonal naive | 42.63  | 79.29  | 0.798 | 0.813               |
| Mean                  | 50.32  | 74.91  | 0.942 | 0.941               |
| Daily seasonal naive  | 86.96  | 129.23 | 1.628 | 0.904               |
| Naive                 | 250.64 | 258.82 | 4.692 | 1.601               |
| Drift                 | 266.37 | 274.61 | 4.986 | 1.606               |

The result is striking and was not anticipated. **SARIMAX scores 0.710 from a single origin against 0.708 from fourteen** a difference of 0.002, far below anything the significance testing can resolve. The weekly seasonal naive is marginally *better* without reissuing, at 0.798 against 0.813.

## 10. Discussion and limitations

Several core constraints limit these results. Operating system variations shift model MASE by  $\sim 0.01$ , and the single 14-day test window (336 correlated points) lacks statistical power: the minimum detectable MASE difference is  $\sim 0.20$ , far larger than any observed gap among the top models. Model specifications were also constrained: SARIMAX omitted weekly seasonality despite its strong empirical presence, and 24-hour origin spacing confounded lead time with time of day (a 23- or 25-hour step would fix this). Finally, operational weather forecasts were unavailable, and pipeline debugging revealed key fit issues in both SARIMAX and gradient boosting.

## 11. Conclusion and deployment recommendation

**1. Which benchmark is strongest, and what does it reveal?** The weekly seasonal naive (MASE 0.813), ahead of the daily variant (0.904). The household's routine is day-specific rather than weekday-weekend: Monday consumption exceeds Tuesday by 28 per cent while the weekend premium is only 4 per cent. A rule looking back one week preserves this; one looking back one day averages it away. That the mean forecast ranks second on RMSE while sixth on MAE further reveals a load dominated by unpredictable spikes, where refusing to guess their timing is a viable strategy under squared-error loss.

**2. Does SARIMAX improve on the strongest benchmark?** By point estimate, yes: 0.708 conditional and 0.721 operational against 0.813, an 11 to 13 per cent improvement. Statistically, no: the Diebold-Mariano test returns  $p = 0.34$ . On this evidence the improvement is unproven. Section 9.4 sharpens the point: SARIMAX scores 0.710 from a single origin against 0.708 from fourteen, so whatever it is doing does not depend on the observations the rolling protocol supplies.

**3. Do lag, rolling, time, sensor and weather features improve the feature-based model?** The calendar features do; the rest do not. Calendar features carry 90.7 per cent of permutation importance, lagged exogenous channels 5.5 per cent, and target lags and rolling statistics 3.9 per cent between them. At a 24-hour horizon the series' own recent history is close to uninformative, and the model reduces to an estimated daily profile one that needs just 25 boosting iterations to express. The model does, however, achieve the best RMSE in the study (65.17), handling consumption spikes marginally better than SARIMAX.

**4. Does the foundation model outperform the simpler models?** On MASE, yes 0.678 against 0.708 for conditional SARIMAX, achieved with no covariates and no training on this dwelling. On RMSE, emphatically not: 74.40 against 66.41, essentially tied with a flat mean forecast. Its bias of  $-29.87$  is six times any other

model's. As Section 8.3 sets out, this pattern follows from summarising a predictive distribution by its median on a right-skewed target, so the MASE advantage partly measures a difference in summary statistic rather than in forecasting ability. The defensible conclusion is that zero-shot transfer performs **comparably** to models fitted on four months of this dwelling's history which is notable but not better.

**5. Which covariates would genuinely be known at the forecast origin?** Calendar features unconditionally. Indoor sensors only at lags of at least the horizon, at which point they contribute 5.5 per cent of importance jointly across eighteen channels — and they should be lagged in any case, since indoor temperature and humidity respond to appliance activity rather than driving it. Weather only via a third-party forecast; the conditional and operational SARIMAX variants bound the cost of not having one at 1.9 per cent of MASE, which the Diebold–Mariano test cannot distinguish from zero ( $p = 0.136$ ). Replacing realised weather with naive persistence costs nothing measurable.

## **6. Which model for a practical smart-home system?**

Accuracy at the operational horizon. The honest figures are 0.721 for operational SARIMAX, 0.711 for the feature model and 0.813 for the weekly seasonal naive a spread no test in this study distinguishes from noise. Section 9.2 shows SARIMAX's advantage concentrating at short leads the 24-hour task does not reward, and Section 9.4 shows it performing identically without any reissuing at all. Marginal gain over the benchmark. At most 13 per cent, statistically unconfirmed ( $p = 0.34$ ), on one household over one fortnight. Operational cost. SARIMAX requires periodic re-estimation and failed to converge under default settings. The gradient boosting model requires a feature pipeline whose correctness must be maintained. The weekly seasonal naive requires a week of stored observations and no maintenance at all. Failure behaviour. A seasonal naive rule copies a real past observation and cannot produce an impossible value. A gradient boosting model extrapolating outside its training range can.

## **References**

- Ansari, A.F., Stella, L., Turkmen, C., Zhang, X., Mercado, P., Shen, H., Shchur, O., Rangapuram, S.S., Arango, S.P., Kapoor, S., Zschiegner, J., Maddix, D.C., Mahoney, M.W., Torkkola, K., Wilson, A.G., Bohlke-Schneider, M. and Wang, Y. (2024) 'Chronos: Learning the Language of Time Series', *Transactions on Machine Learning Research*.
- Candanedo, L.M., Feldheim, V. and Deramaix, D. (2017) 'Data driven prediction models of energy use of appliances in a low-energy house', *Energy and Buildings*, 140, pp. 81–97.
- Diebold, F.X. and Mariano, R.S. (1995) 'Comparing predictive accuracy', *Journal of Business and Economic Statistics*, 13(3), pp. 253–263.
- Hyndman, R.J. and Athanasopoulos, G. (2021) *Forecasting: Principles and Practice*. 3rd edn. Melbourne: OTexts.
- Hyndman, R.J. and Koehler, A.B. (2006) 'Another look at measures of forecast accuracy', *International Journal of Forecasting*, 22(4), pp. 679–688.
- Makridakis, S., Spiliotis, E. and Assimakopoulos, V. (2018) 'The M4 Competition: Results, findings, conclusion and way forward', *International Journal of Forecasting*, 34(4), pp. 802–808.
- Newey, W.K. and West, K.D. (1987) 'A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix', *Econometrica*, 55(3), pp. 703–708.
- Tashman, L.J. (2000) 'Out-of-sample tests of forecasting accuracy: an analysis and review', *International Journal of Forecasting*, 16(4), pp. 437–450.
- Wang, X., Smith, K.A. and Hyndman, R.J. (2006) 'Characteristic-based clustering for time series data', *Data Mining and Knowledge Discovery*, 13(3), pp. 335–364.